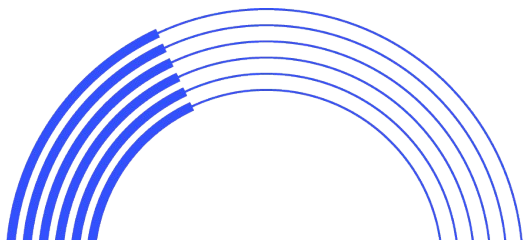



Distributed Database

2021/06/17




Agenda

1. **The Evolution of Distributed Database**
2. **TiDB Architecture**
3. **How to Participate in This Area**



The Evolution of Distributed Database



Overview

1. **Sharding Middleware**
2. **NoSQL**
3. **NewSQL/HTAP**
4. **Cloud Native**

Sharding Middleware

To solve the capacity bottleneck in single node:

- combine multi database nodes together into a logical database
- route queries to the correct node

Common problems:

- application developers need design the schema and queries carefully, resharding is painful
- lacking some features, for example, cross node join (depends on the implementation)

Typical systems: Vitess, MyCat, DRDS, ShardingSphere, etc.

NoSQL

To solve the scalability problem and achieve high availability:

- no-relational data model
- low latency, high throughput, eventually consistency

Common problems:

- No SQL API, not ideal for complex applications
- No strong transactional semantics, not ideal for some mission-critical applications

Typical systems: BigTable, HBase, Cassandra, MongoDB, DynamoDB

NewSQL

To provide scalability, flexibility, high availability, ACID transaction, and SQL support:

- Deliver the scalability and flexibility promised by NoSQL systems
- Providing ACID transaction and SQL support

Common problems:

- Transaction latency is High without fine-grained optimization

Typical systems: Spanner, F1, CockroachDB, TiDB


Some of them are working on HTAP functionality

Cloud Native


To provide scalability, flexibility, high availability, ACID transaction, and SQL support:

- Building database components based on cloud services
- Scale storage and compute independently

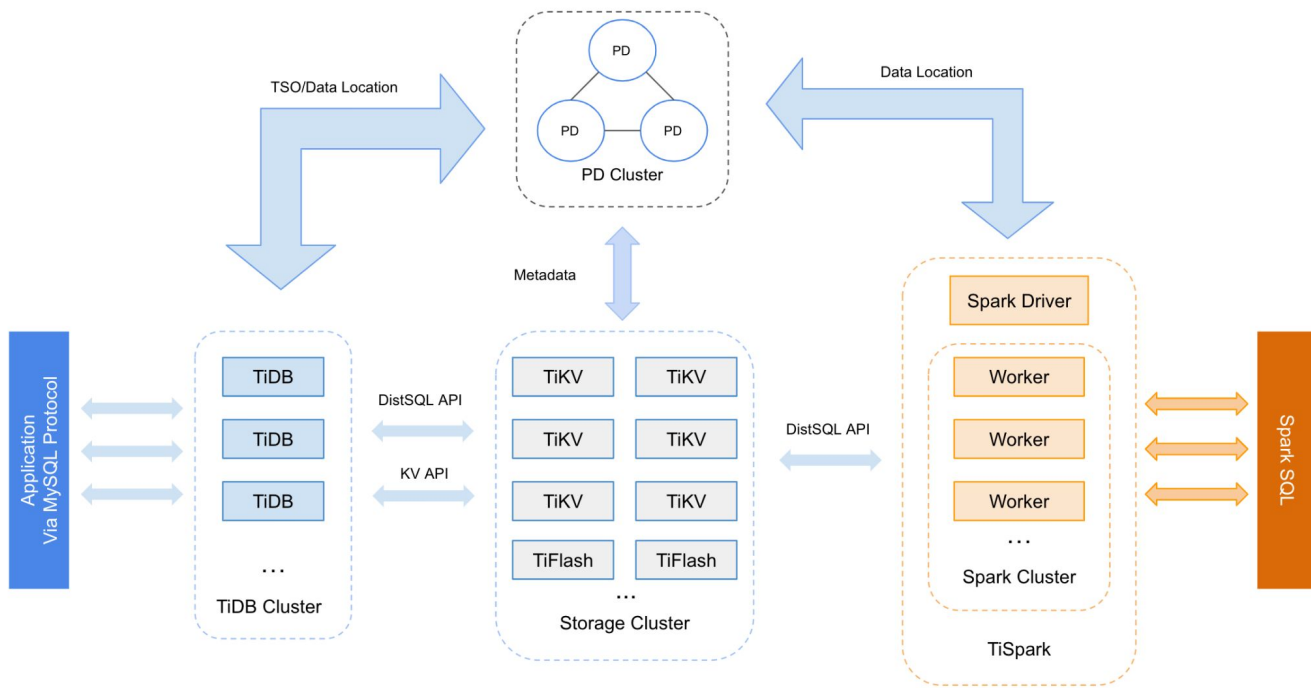
Typical systems: Snowflake, Aurora, Cloud Spanner



TiDB Architecture



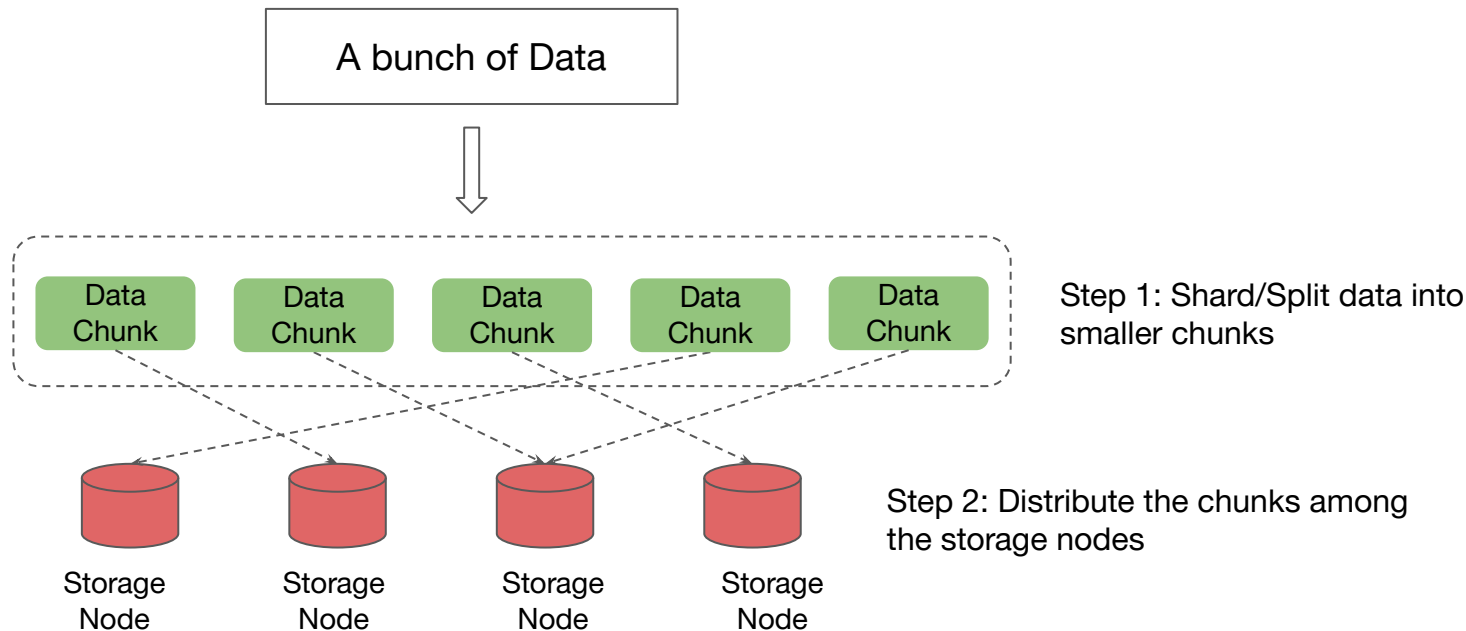
Overview



TiDB Design Goals

- **Horizontal Scalability**
 - Scale-out process should be transparent to the application layer
- **High Availability**
 - Self-healing thanks to modern consensus algorithm like Raft/Multi-Paxos
- **ACID Compliance**
 - Distributed cross-row transaction support
- **MySQL Protocol & Dialect**
 - MySQL syntax & ecosystem tools


To achieve Horizontal Scalability



How to decide chunk size?

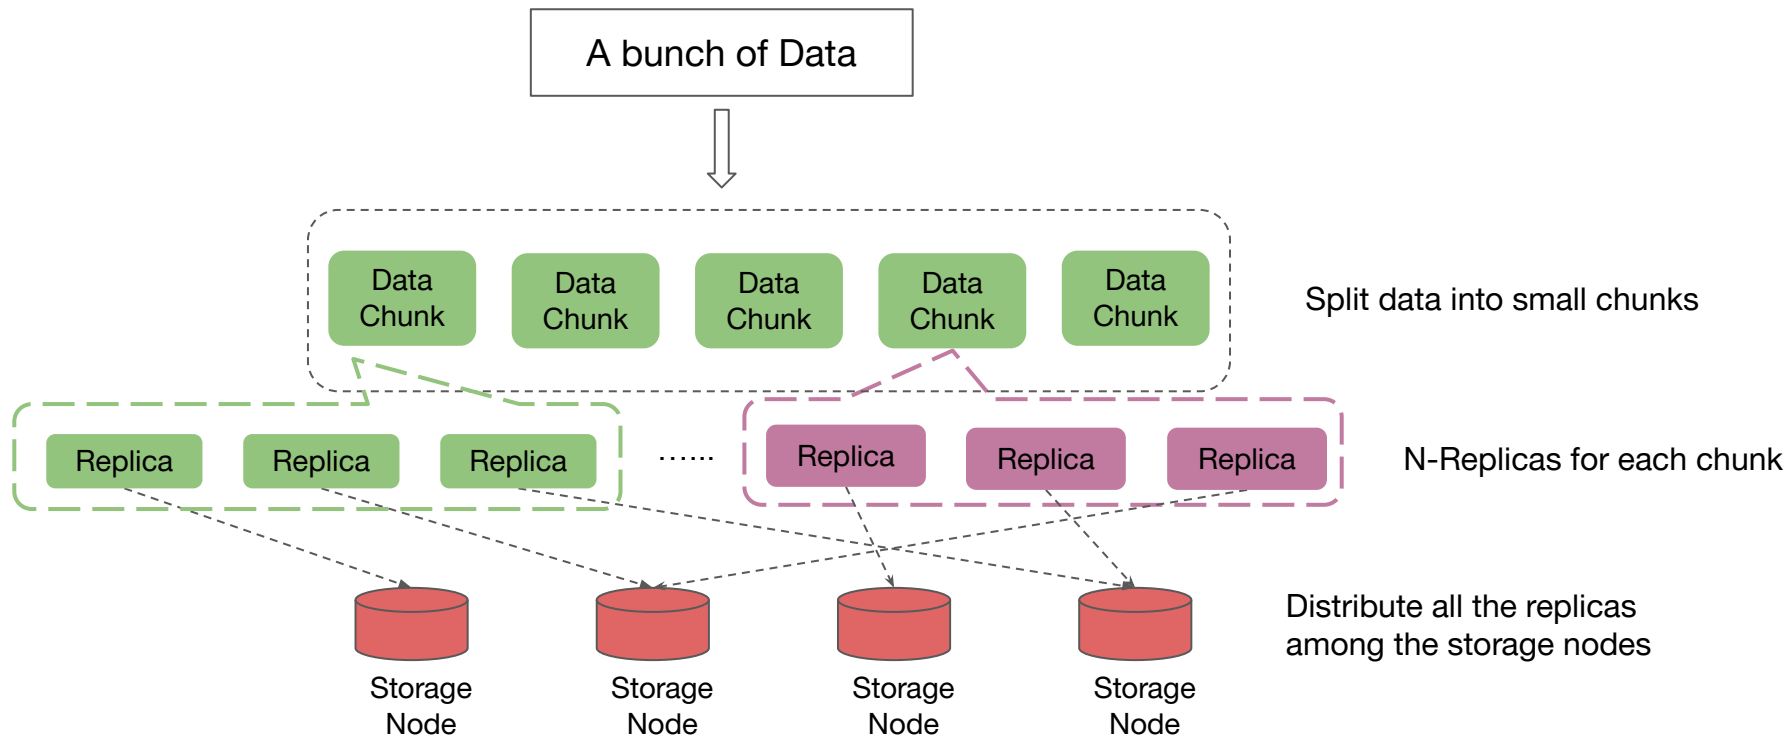
How to do balance load?

To achieve High Availability

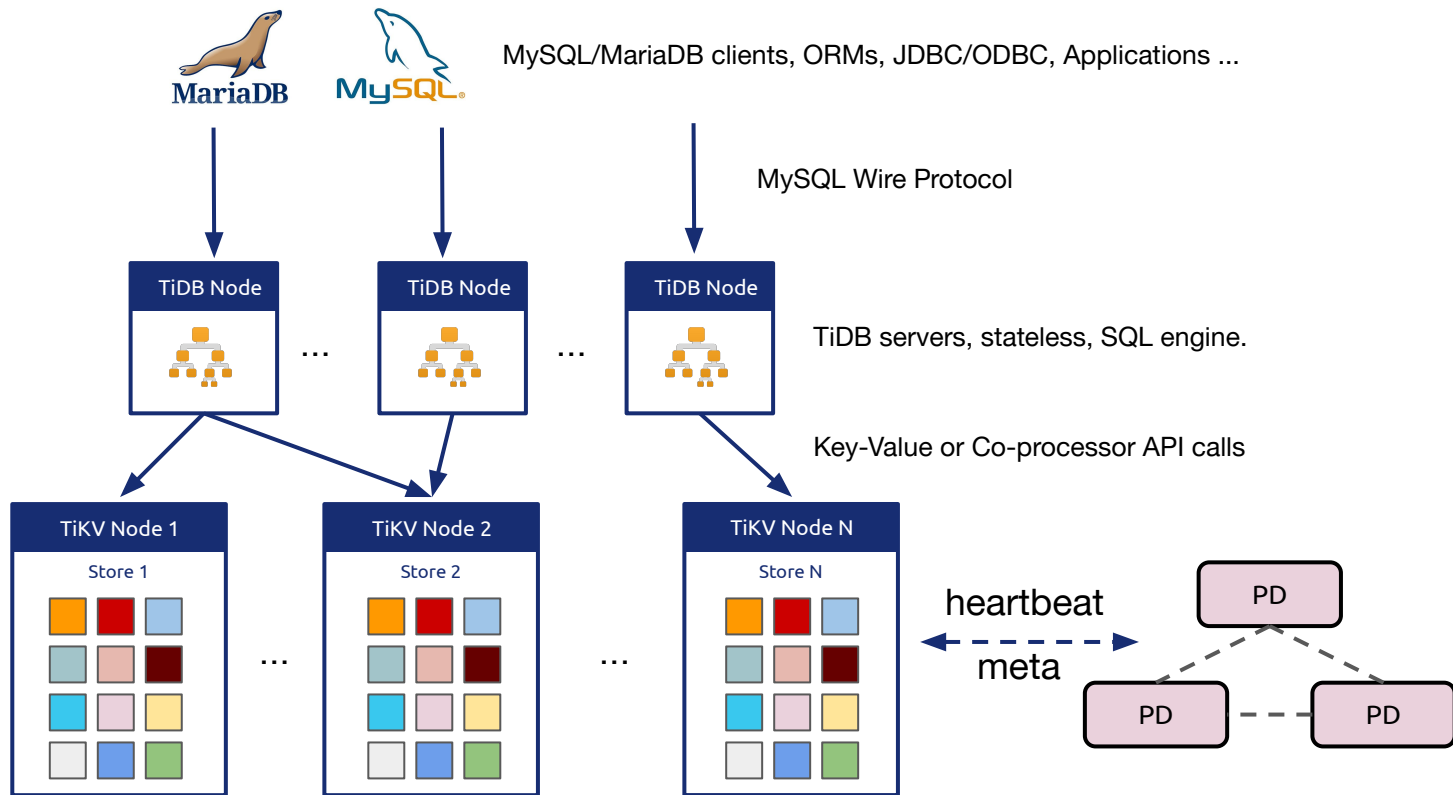


How to make data consistent among replicas?
What will happen if one replica is gone?

Combine Scalability and HA together



TiDB Architecture



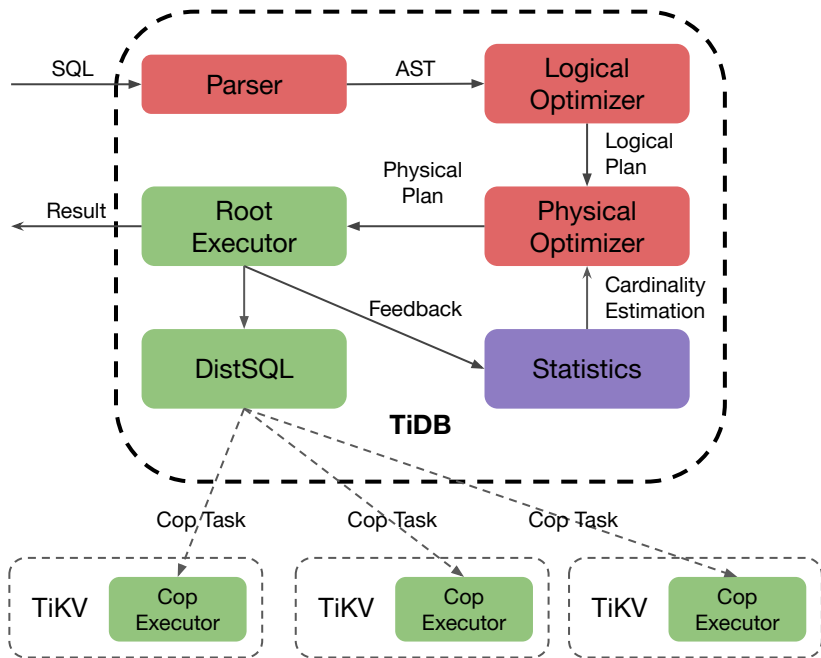
Query Execution

Query Optimization

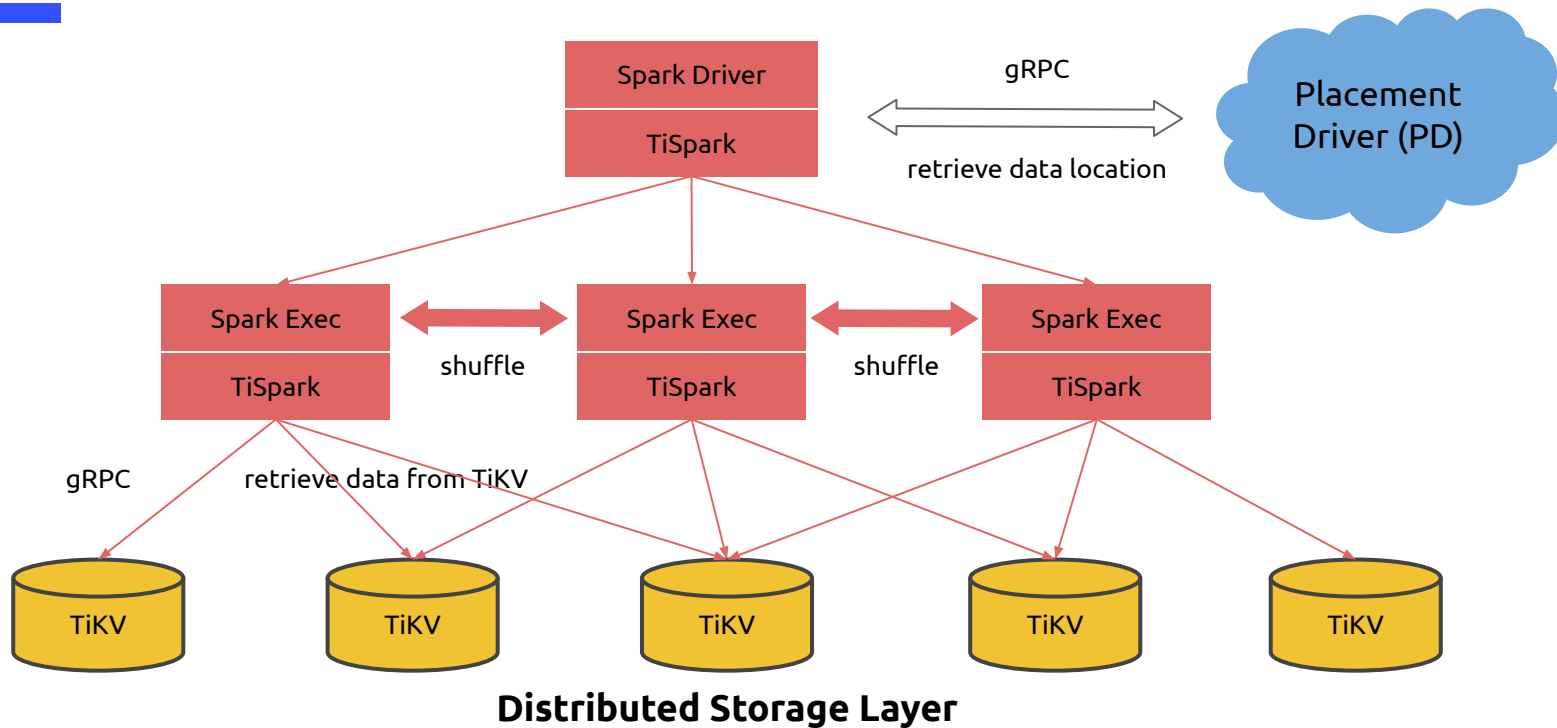
- Logical Optimizer
- Physical Optimizer
- Statistics

Query Execution

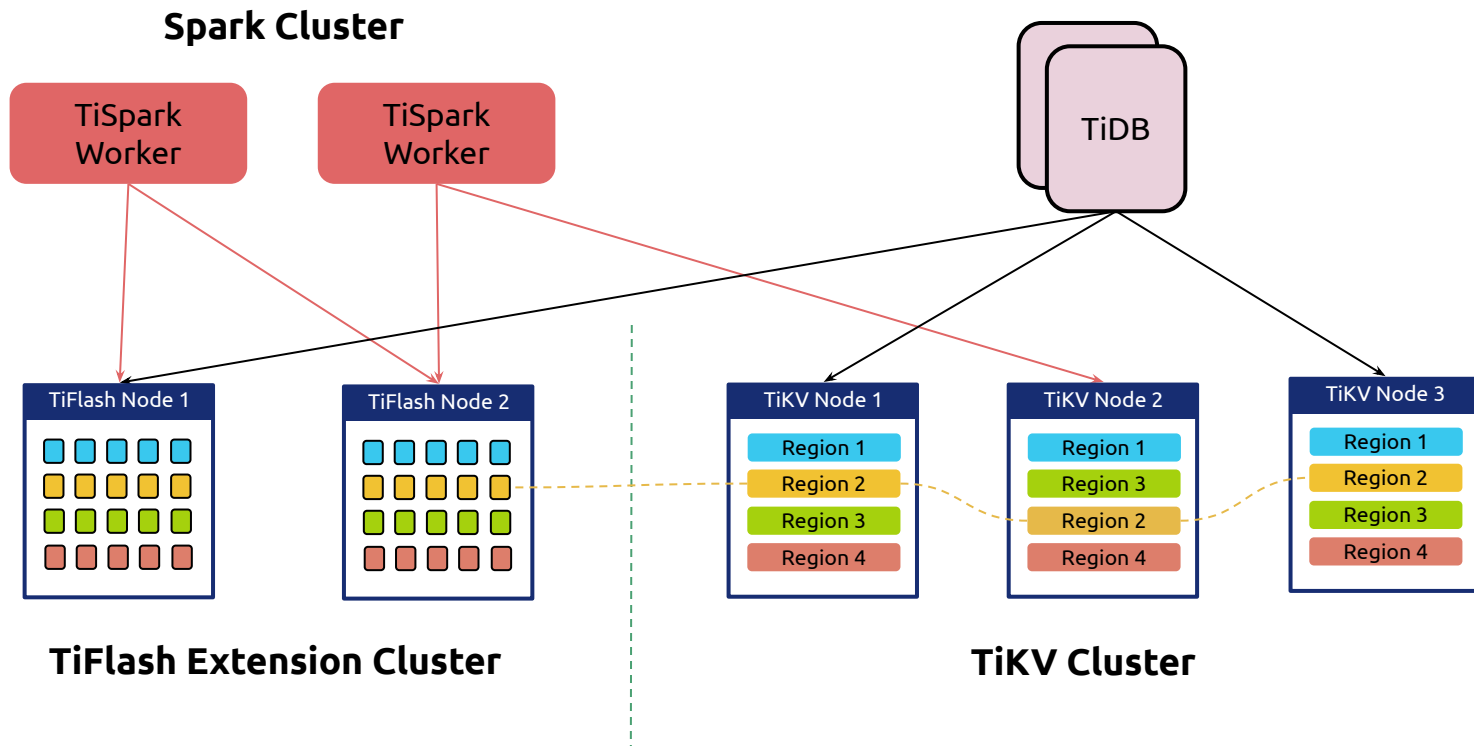
- Root Executor
- DistSQL
- Coprocessor Executor



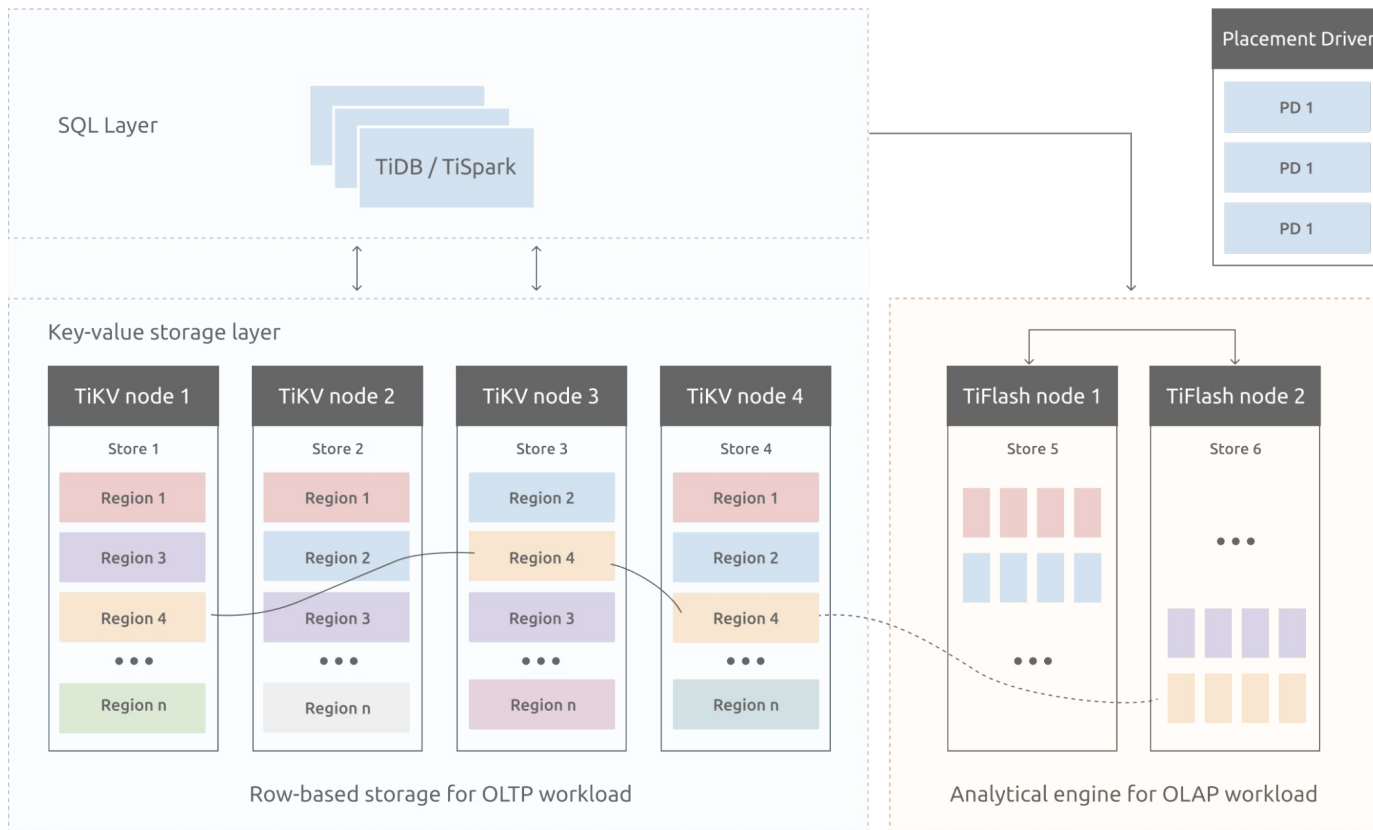
TiSpark



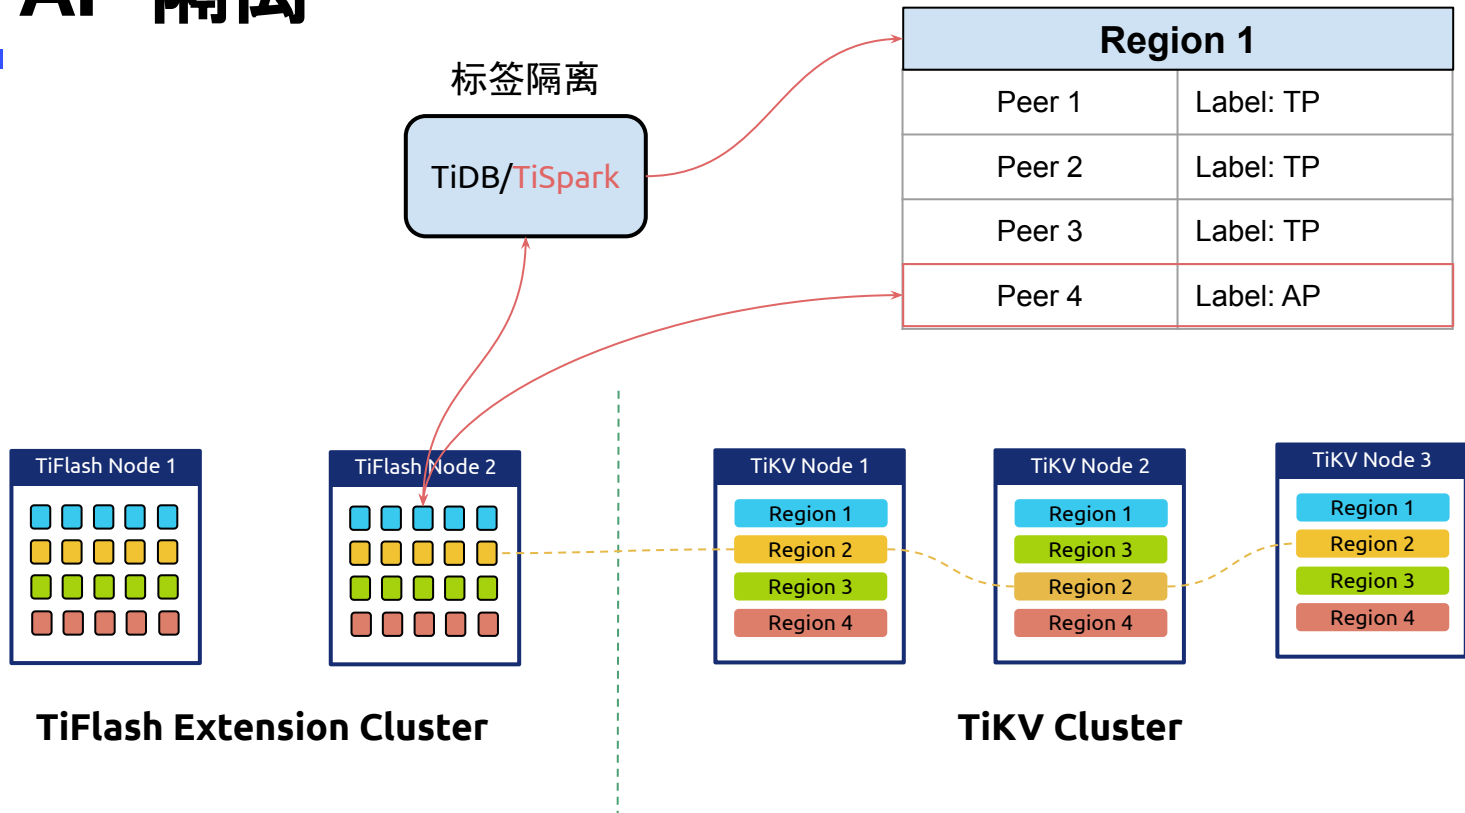
TiFlash 4.0



TiFlash 5.0 - MPP Support

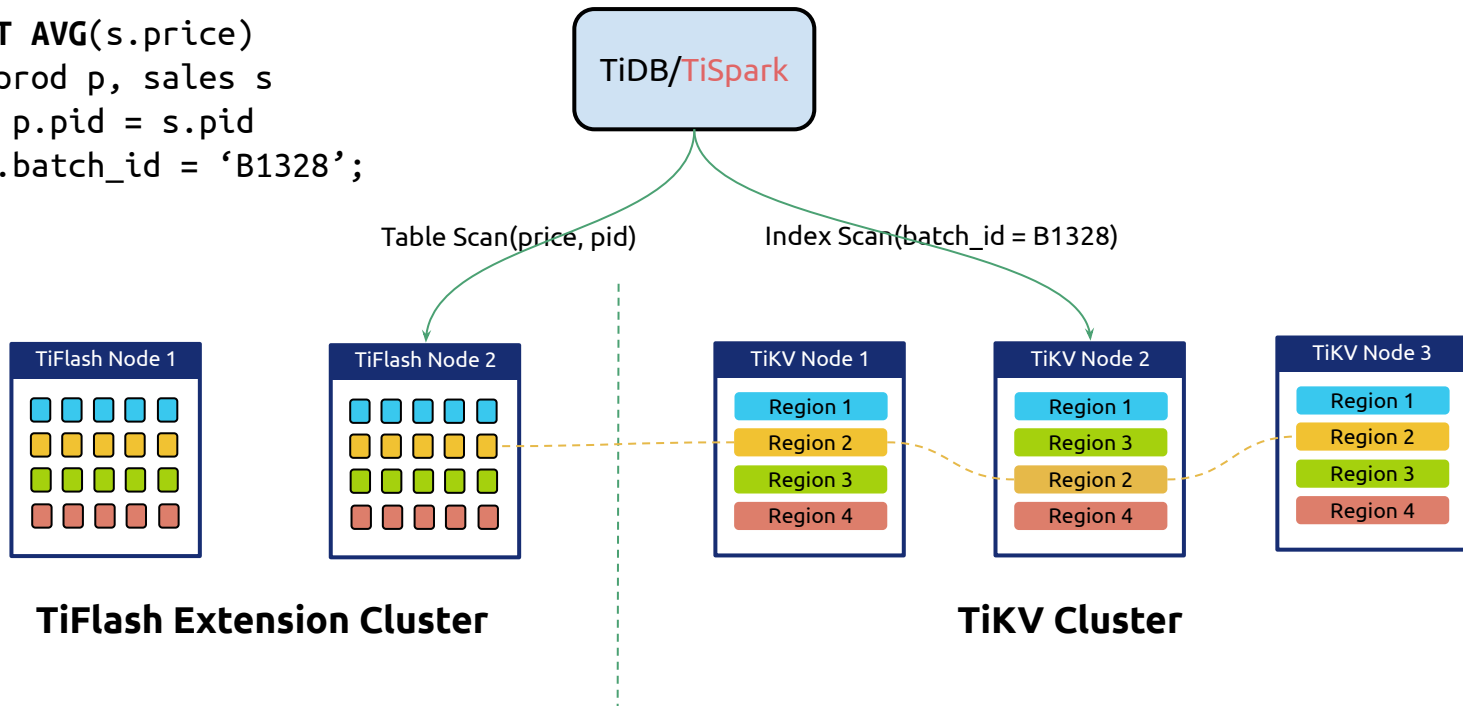



TP、AP 隔离




TP、AP 融合

```
SELECT AVG(s.price)
FROM prod p, sales s
WHERE p.pid = s.pid
AND p.batch_id = 'B1328';
```





How to Participate in This Area

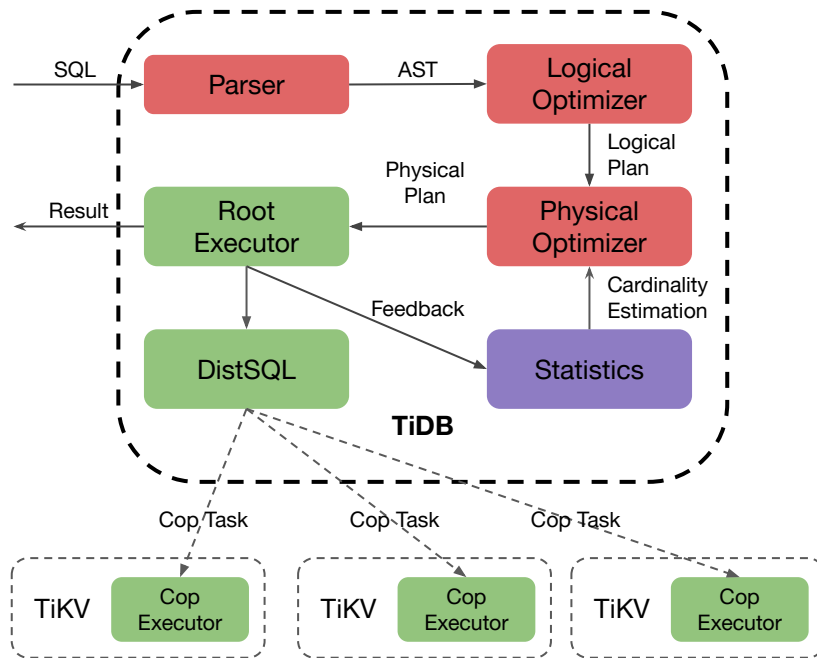


Suggestions

从一个小模块开始, 成为细分领域的专家

逐渐扩展多个小模块, 成为它们的细分领域专家

Suggestions



Suggestions

